

PDE-based Pricing in C++

An explicit finite-difference solver for Black–Scholes with dividends

Programming Project Report

Sacha Faye Mathys Gouin Samuel Macé

Repository (public):

https://github.com/Samgit0532/Pricing_par_EDP

Abstract. We implement a modular C++ pricing engine for several derivative products by numerically solving the Black–Scholes–Merton PDE with continuous dividends using an explicit finite-difference scheme. The code supports European and American derivatives (via early-exercise projection), returns price and Greeks (Delta, Gamma), and includes a test suite validating basic financial identities and numerical sanity checks (put–call parity, forward closed-form, dominance of American options, and composite product consistency). A central engineering contribution is an automatic grid construction driven by a user-chosen spatial resolution parameter `rel_dS`, such that $\Delta S = \text{rel_dS} \cdot S_0$, designed to preserve accuracy near the spot while keeping runtimes reasonable.

Contents

1 Project architecture (high-level overview)	2
1.1 How modules interact	2
1.2 What the program takes as inputs	2
2 Model: Black–Scholes–Merton with dividends	3
2.1 Risk-neutral dynamics	3
2.2 Pricing PDE	3
2.3 Implementation in our code	3
3 Solver: explicit finite differences (<code>solvers/ExplicitFdSolver.hpp</code>, <code>solvers/Solver.cpp</code>)	3
3.1 Terminal condition	3
3.2 Explicit backward time stepping	3
3.3 Boundary conditions	4
3.4 American early exercise	4
3.5 Outputs of the solver	4
4 Implemented products (detailed)	4
4.1 European Call (<code>products/EuropeanCall.hpp</code>)	5
4.2 European Put (<code>products/EuropeanPut.hpp</code>)	5
4.3 American Put (<code>products/AmericanPut.hpp</code>)	5
4.4 American Call (<code>products/AmericanCall.hpp</code>)	5
4.5 Forward/Future-like contract (<code>products/Future.hpp</code>)	5
4.6 Bull Call Spread (<code>products/BullCallSpread.hpp</code>)	6
4.7 Bear Put Spread (<code>products/BearPutSpread.hpp</code>)	6
4.8 Straddle (<code>products/Straddle.hpp</code>)	6
5 Finite-difference grid and interpolation	6
5.1 Grid definition (<code>grid/FdGrid.hpp</code>)	6
5.2 Interpolation used for the final price	7
6 Automatic grid parameter selection: stability vs accuracy vs speed	7
6.1 Domain truncation: choosing S_{\max} via a lognormal quantile	7
6.2 Space resolution: choosing N_s from a relative mesh size	7
6.3 Time resolution: explicit stability constraint and its implications	8
6.4 Main criticism of the approach	8
7 Greeks: formulas and numerical computation	8
7.1 Definitions (continuous-time)	8
7.2 Discrete formulas used in our code	9
8 Test suite: consistency checks	9

1 Project architecture (high-level overview)

Our codebase is organized as follows (rooted in `src/`):

```

src/
    main.cpp                  (interactive application)
    tests/
        TestPricing.cpp      (automated consistency tests)
    model/
        BlackScholesModel.hpp
    grid/
        FdGrid.hpp
        GridParameters.hpp
    products/
        InterfaceProducts.hpp
        EuropeanCall.hpp     EuropeanPut.hpp
        AmericanCall.hpp    AmericanPut.hpp
        Future.hpp
        BullCallSpread.hpp   BearPutSpread.hpp
        Straddle.hpp
    solvers/
        ExplicitFdSolver.hpp
        Solver.cpp           (implementation of ExplicitFdSolver::price)

```

1.1 How modules interact

The core pricing pipeline is always the same:

1. We define market parameters in `model/BlackScholesModel.hpp`.
2. We instantiate a product (payoff + boundary conditions) from `products/`.
3. We build a finite-difference grid using `grid/GridParameters.hpp`, which returns an `FdGrid` from `grid/FdGrid.hpp`.
4. We call `ExplicitFdSolver::price` (declared in `solvers/ExplicitFdSolver.hpp`, implemented in `solvers/Solver.cpp`).
5. The solver returns price + Greeks (Delta, Gamma) and the entire solution slice at $t = 0$.

1.2 What the program takes as inputs

At runtime (interactive application in `src/main.cpp`), the user provides:

- Market inputs: S_0 , r , σ , q .
- Contract inputs: maturity T , strike(s) depending on product.
- Numerical input: a spatial resolution parameter `rel_dS` such that $\Delta S = \text{rel_dS} \cdot S_0$.

The test program (`src/tests/TestPricing.cpp`) uses fixed parameter sets and verifies financial identities and numerical properties automatically.

2 Model: Black–Scholes–Merton with dividends

2.1 Risk-neutral dynamics

We assume the underlying price S_t follows (under the risk-neutral measure):

$$dS_t = (r - q)S_t dt + \sigma S_t dW_t,$$

where r is the risk-free rate, σ the volatility, and q the continuous dividend yield.

2.2 Pricing PDE

For a derivative value $V(t, S)$, the Black–Scholes–Merton PDE is:

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + (r - q)S \frac{\partial V}{\partial S} - rV = 0.$$

The PDE is solved backward in time with:

- terminal condition (payoff): $V(T, S) = \Phi(S)$,
- boundary conditions as $S \rightarrow S_{\min}$ and $S \rightarrow S_{\max}$.

2.3 Implementation in our code

The model parameters are stored in `model/BlackScholesModel.hpp`:

$$(r, \sigma, q),$$

and used both:

- in `solvers/Solver.cpp` to build finite-difference coefficients (drift $r - q$, discount rate r),
- in product boundary conditions (e.g. call boundary uses $e^{-q(T-t)}$ and $e^{-r(T-t)}$).

3 Solver: explicit finite differences (`solvers/ExplicitFdSolver.hpp`, `solvers/Solver.cpp`)

3.1 Terminal condition

For each product, we initialize:

$$V_i^{N_t} = \Phi(S_i),$$

where Φ is the payoff, implemented as `option.payoff(S[i])`.

3.2 Explicit backward time stepping

Let $V_i^n \approx V(t_n, S_i)$. For each time step $n = N_t - 1, \dots, 0$, we compute for interior nodes $i = 1, \dots, N_s - 1$:

$$V_i^n = A_i V_{i-1}^{n+1} + B_i V_i^{n+1} + C_i V_{i+1}^{n+1},$$

with:

$$\begin{aligned} A_i &= \frac{\Delta t}{2} \left(\frac{\sigma^2 S_i^2}{\Delta S^2} - \frac{(r - q)S_i}{\Delta S} \right), \\ B_i &= 1 - \Delta t \left(\frac{\sigma^2 S_i^2}{\Delta S^2} + r \right), \\ C_i &= \frac{\Delta t}{2} \left(\frac{\sigma^2 S_i^2}{\Delta S^2} + \frac{(r - q)S_i}{\Delta S} \right). \end{aligned}$$

This is implemented in `solvers/Solver.cpp` with the variables A, B, C.

3.3 Boundary conditions

At each time layer t_n , boundaries are imposed through the product interface:

$$V_0^n = \text{option.leftBoundary}(t_n), \quad V_{N_s}^n = \text{option.rightBoundary}(t_n, S_{\max}).$$

This design makes it straightforward to add products: each product specifies its asymptotic behaviour.

3.4 American early exercise

For American options, after computing the continuation value, we project:

$$V_i^n \leftarrow \max(V_i^n, \Phi(S_i)),$$

implemented via: `if(option.isAmerican()) Vnew[i]=max(Vnew[i], option.earlyExerciseValue(Si));`. This is the standard approach for PDE-based American pricing.

3.5 Outputs of the solver

The solver returns a `Result` struct (see `solvers/ExplicitFdSolver.hpp`) containing:

- `V0`: the full vector $\{V(0, S_i)\}_{i=0}^{N_s}$,
- `price`: interpolated $V(0, S_0)$,
- `delta`, `gamma`: computed at S_0 using finite differences.

4 Implemented products (detailed)

Before discussing grid automation and numerical tuning, we present the set of derivatives supported by our engine. Each product implements the common interface in `products/InterfaceProducts.hpp`:

- `maturity()` and `strike()` (or a representative strike),
- `payoff(S)` (terminal condition),
- `leftBoundary(t)` and `rightBoundary(t, Smax)`,
- optional American hooks: `isAmerican()`, `earlyExerciseValue(S)`.

4.1 European Call (`products/EuropeanCall.hpp`)

Inputs: strike K , maturity T , model (r, σ, q) .

Payoff:

$$\Phi(S) = \max(S - K, 0).$$

Boundaries:

$$V(t, 0) = 0, \quad V(t, S_{\max}) \approx S_{\max}e^{-q(T-t)} - Ke^{-r(T-t)}.$$

4.2 European Put (`products/EuropeanPut.hpp`)

Inputs: $K, T, (r, \sigma, q)$.

Payoff:

$$\Phi(S) = \max(K - S, 0).$$

Boundaries:

$$V(t, 0) \approx Ke^{-r(T-t)}, \quad V(t, S_{\max}) \approx 0.$$

4.3 American Put (`products/AmericanPut.hpp`)

Payoff: $\max(K - S, 0)$.

Early exercise: enabled via `isAmerican()=true`, with:

$$V(t, S) \geq \max(K - S, 0).$$

Left boundary choice: we set $V(t, 0) = K$. This reflects the fact that when $S = 0$, immediate exercise yields K and is optimal.

4.4 American Call (`products/AmericanCall.hpp`)

Payoff: $\max(S - K, 0)$.

Early exercise: enabled. In Black–Scholes, early exercise for calls is:

- never optimal when $q = 0$ (American call equals European call),
- potentially optimal when $q > 0$ (dividends can make early exercise attractive).

Boundary at S_{\max} : same asymptotic as European call:

$$V(t, S_{\max}) \approx S_{\max}e^{-q(T-t)} - Ke^{-r(T-t)}.$$

4.5 Forward/Future-like contract (`products/Future.hpp`)

Inputs: delivery price K , maturity T .

Payoff:

$$\Phi(S) = S - K.$$

This product admits a closed-form price:

$$V(0, S_0) = S_0e^{-qT} - Ke^{-rT}.$$

Boundaries:

$$V(t, 0) \approx -Ke^{-r(T-t)}, \quad V(t, S_{\max}) \approx S_{\max}e^{-q(T-t)} - Ke^{-r(T-t)}.$$

4.6 Bull Call Spread (products/BullCallSpread.hpp)

Inputs: strikes $K_1 < K_2$, maturity T .

Payoff:

$$\Phi(S) = \max(S - K_1, 0) - \max(S - K_2, 0).$$

This payoff is bounded between 0 and $(K_2 - K_1)$.

Right boundary: for very large S , the payoff saturates to $K_2 - K_1$, hence:

$$V(t, S_{\max}) \approx (K_2 - K_1)e^{-r(T-t)}.$$

Left boundary: $V(t, 0) = 0$.

4.7 Bear Put Spread (products/BearPutSpread.hpp)

Inputs: $K_1 < K_2$, T .

Payoff:

$$\Phi(S) = \max(K_2 - S, 0) - \max(K_1 - S, 0),$$

also bounded between 0 and $(K_2 - K_1)$.

Left boundary: for $S = 0$, payoff saturates to $K_2 - K_1$, so:

$$V(t, 0) \approx (K_2 - K_1)e^{-r(T-t)}.$$

Right boundary: $V(t, S_{\max}) \approx 0$.

4.8 Straddle (products/Straddle.hpp)

Inputs: strike K , maturity T .

Payoff:

$$\Phi(S) = |S - K| = \max(S - K, 0) + \max(K - S, 0).$$

A key identity is:

$$\text{Straddle} = \text{Call}(K) + \text{Put}(K).$$

Boundaries:

$$V(t, 0) \approx Ke^{-r(T-t)}, \quad V(t, S_{\max}) \approx S_{\max}e^{-q(T-t)} - Ke^{-r(T-t)}.$$

5 Finite-difference grid and interpolation

5.1 Grid definition (grid/FdGrid.hpp)

We discretize:

$$t_n = n\Delta t, \quad n = 0, \dots, N_t, \quad S_i = S_{\min} + i\Delta S, \quad i = 0, \dots, N_s.$$

The `FdGrid` class stores:

$$T, S_{\min}, S_{\max}, N_t, N_s, \Delta t, \Delta S,$$

and exposes `timeGrid()`, `priceGrid()`, plus a safe linear interpolation routine `interpolate(V, S0)` to evaluate $V(0, S_0)$ even when S_0 is not exactly a grid node.

5.2 Interpolation used for the final price

After solving for the vector $V(0, S_i)$ on the grid, we compute:

$$V(0, S_0) \approx (1 - w)V(0, S_i) + wV(0, S_{i+1}), \quad w = \frac{S_0 - S_i}{S_{i+1} - S_i},$$

with (S_i, S_{i+1}) such that $S_i \leq S_0 \leq S_{i+1}$. This is implemented in `FdGrid::interpolate`.

6 Automatic grid parameter selection: stability vs accuracy vs speed

A key engineering challenge of this project was to select grid parameters that simultaneously ensure:

- **stability** (explicit scheme \Rightarrow CFL-type constraint),
- **accuracy** (small discretization error near S_0 and around the strike),
- **efficiency** (runtime and memory remain reasonable for interactive use).

In practice, this was the main difficulty: a naive grid choice can easily produce a solver that is either (i) unstable and diverges, or (ii) stable but too slow, or (iii) fast but inaccurate because the spatial mesh becomes too coarse.

6.1 Domain truncation: choosing S_{\max} via a lognormal quantile

Under Black–Scholes, $\log S_T$ is Gaussian:

$$\log S_T \sim \mathcal{N}\left(\log S_0 + (r - q - \frac{1}{2}\sigma^2)T, \sigma^2 T\right).$$

We choose the upper boundary as a high quantile of the lognormal distribution:

$$S_{\max} = S_0 \exp\left((r - q - \frac{1}{2}\sigma^2)T + z\sigma\sqrt{T}\right),$$

with a fixed safety level z (e.g. $z = 5$). The motivation is to make the probability of $S_T > S_{\max}$ negligible, so that boundary conditions do not contaminate the solution in the region of interest (near S_0). We set $S_{\min} = 0$ for all products.

6.2 Space resolution: choosing N_s from a relative mesh size

A first implementation used fixed values of N_s (e.g. 250/400/800). This worked on some parameter sets but degraded when S_{\max} increased (large T or large σ). Since

$$\Delta S = \frac{S_{\max} - S_{\min}}{N_s},$$

a larger S_{\max} mechanically increases ΔS and reduces accuracy around S_0 .

To fix this, our final design makes the spatial step proportional to the spot:

$$\Delta S = \text{rel_dS} \cdot S_0,$$

where `rel_dS` is chosen by the user at runtime (typical values: 0.004 fast, 0.002 balanced, 0.001 accurate). We then set:

$$N_s = \left\lceil \frac{S_{\max} - S_{\min}}{\Delta S} \right\rceil.$$

This ensures that local resolution near S_0 remains comparable across parameter sets, even when S_{\max} varies substantially.

6.3 Time resolution: explicit stability constraint and its implications

The explicit finite-difference scheme is only conditionally stable, so Δt must satisfy a CFL-type bound. A conservative practical constraint is:

$$\Delta t \lesssim \frac{1}{\sigma^2 S_{\max}^2 / \Delta S^2 + r}.$$

We choose Δt using this worst-case bound (with a safety factor) and set:

$$N_t = \left\lceil \frac{T}{\Delta t} \right\rceil.$$

6.4 Main criticism of the approach

The above procedure works reliably for many cases, but it also highlights the main limitation of our numerical method:

- Because we use an **explicit** scheme, stability can force N_t to become very large when σ or S_{\max} is large, which may make the solver slow.
- Using a uniform grid in S can still be inefficient: accuracy is needed mainly near S_0 (and around strikes), while far-away regions consume grid points but contribute little to the price at S_0 .

These limitations are well-known in PDE pricing: implicit schemes (Backward Euler / Crank–Nicolson) remove the strict stability constraint, and non-uniform grids (or a change of variable to $\log S$) concentrate points where they matter most. We chose the explicit scheme for simplicity and transparency, but grid automation was crucial to obtain a solver that is both reasonably fast and accurate in practice.

7 Greeks: formulas and numerical computation

7.1 Definitions (continuous-time)

For a derivative price $V(0, S_0; r, \sigma, q, T, \dots)$, the standard Greeks are:

$$\begin{aligned} \Delta &= \frac{\partial V}{\partial S_0}, & \Gamma &= \frac{\partial^2 V}{\partial S_0^2}, & \Theta &= \frac{\partial V}{\partial t}, \\ \nu \text{ (Vega)} &= \frac{\partial V}{\partial \sigma}, & \rho &= \frac{\partial V}{\partial r}. \end{aligned}$$

In this project, we compute **Delta** and **Gamma** directly from the PDE solution on the S -grid. The other Greeks (Theta, Vega, Rho) can be computed by bump-and-revalue, but we chose not to implement them to keep the engine focused.

7.2 Discrete formulas used in our code

Let $V_i \approx V(0, S_i)$ and choose i^* such that S_{i^*} is the closest grid node to S_0 (while avoiding boundaries). Then:

$$\Delta(0, S_0) \approx \frac{V_{i^*+1} - V_{i^*-1}}{2\Delta S},$$

$$\Gamma(0, S_0) \approx \frac{V_{i^*+1} - 2V_{i^*} + V_{i^*-1}}{\Delta S^2}.$$

We avoid $i^* = 0$ and $i^* = N_s$ to prevent unreliable boundary derivatives; this is implemented in `solvers/Solver.cpp` by shifting the index away from the boundaries.

8 Test suite: consistency checks

The test executable `bs_tests` runs a suite of checks comparing our numerical outputs to basic financial identities and numerical properties:

- **Put–call parity (European):**

$$C - P = S_0 e^{-qT} - K e^{-rT}.$$

- **Forward closed-form price:**

$$V_{\text{fwd}}(0) = S_0 e^{-qT} - K e^{-rT}.$$

- **American dominance:** $P^{Am} \geq P^{Eu}$ and $C^{Am} \geq C^{Eu}$.
- **Straddle consistency:** Straddle $\approx C + P$.
- **Spread bounds:** spreads are bounded by $(K_2 - K_1)e^{-rT}$.
- **Greek sanity checks:** for forwards, $\Delta \approx e^{-qT}$ and $\Gamma \approx 0$.

These tests were crucial to ensure our code behaves correctly and to detect regressions after refactoring.