

# PDE-based Pricing in C++

An explicit finite-difference solver for Black–Scholes with dividends

---

## Programming Project Report

---

Sacha Faye    Mathys Gouin    Samuel Macé

Repository (public):

[https://github.com/Samgit0532/Pricing\\_par\\_EDP](https://github.com/Samgit0532/Pricing_par_EDP)

**Abstract.** We implement a modular C++ pricing engine for several derivative products by numerically solving the Black–Scholes–Merton PDE with continuous dividends using an explicit finite-difference scheme. The code supports European and American derivatives (via early-exercise projection), returns price and Greeks (Delta, Gamma), and includes a test suite validating basic financial identities and numerical sanity checks (put–call parity, forward closed-form, dominance of American options, and composite product consistency). A central engineering contribution is an automatic grid construction driven by a user-chosen spatial resolution parameter `rel_dS`, such that  $\Delta S = \text{rel\_dS} \cdot S_0$ , designed to preserve accuracy near the spot while keeping runtimes reasonable.

## Contents

<b>1 Project architecture (high-level overview)</b>	<b>3</b>
1.1 How modules interact . . . . .	3
1.2 What the program takes as inputs . . . . .	3
<b>2 Model: Black–Scholes–Merton with dividends</b>	<b>4</b>
2.1 Risk-neutral dynamics . . . . .	4
2.2 Pricing PDE . . . . .	4
2.3 Implementation in our code . . . . .	4
<b>3 Solver: explicit finite differences (<code>solvers/ExplicitFdSolver.hpp</code>, <code>solvers/Solver.cpp</code>)</b>	<b>4</b>
3.1 Terminal condition . . . . .	4
3.2 Explicit backward time stepping . . . . .	4
3.3 Boundary conditions . . . . .	5
3.4 American early exercise . . . . .	5
3.5 Outputs of the solver . . . . .	5
<b>4 Implemented products (detailed)</b>	<b>5</b>
4.1 European Call ( <code>products/EuropeanCall.hpp</code> ) . . . . .	6
4.2 European Put ( <code>products/EuropeanPut.hpp</code> ) . . . . .	6
4.3 American Put ( <code>products/AmericanPut.hpp</code> ) . . . . .	6
4.4 American Call ( <code>products/AmericanCall.hpp</code> ) . . . . .	6
4.5 Forward/Future-like contract ( <code>products/Future.hpp</code> ) . . . . .	6
4.6 Bull Call Spread ( <code>products/BullCallSpread.hpp</code> ) . . . . .	7
4.7 Bear Put Spread ( <code>products/BearPutSpread.hpp</code> ) . . . . .	7
4.8 Straddle ( <code>products/Straddle.hpp</code> ) . . . . .	7
<b>5 Finite-difference grid and interpolation</b>	<b>7</b>
5.1 Grid definition ( <code>grid/FdGrid.hpp</code> ) . . . . .	7
5.2 Interpolation used for the final price . . . . .	8
<b>6 Automatic grid parameter selection: stability vs accuracy vs speed</b>	<b>8</b>
6.1 Domain truncation: choosing $S_{\max}$ via a lognormal quantile . . . . .	8
6.2 Space resolution: choosing $N_s$ from a relative mesh size . . . . .	8
6.3 Time resolution: explicit stability constraint and its implications . . . . .	9
6.4 Main criticism of the approach . . . . .	9
<b>7 Greeks: formulas and numerical computation</b>	<b>9</b>
7.1 Definitions (continuous-time) . . . . .	9
7.2 Discrete formulas used in our code . . . . .	10
<b>8 Test suite: consistency checks</b>	<b>10</b>
<b>9 How to build and run the code</b>	<b>10</b>
9.1 Build and run with CMake . . . . .	10
9.1.1 Requirements . . . . .	10
9.1.2 Build . . . . .	11

9.1.3	Run the interactive application . . . . .	11
9.1.4	Run the test suite . . . . .	11
9.2	Build and run without CMake (direct compilation) . . . . .	11
9.2.1	Requirements . . . . .	11
9.2.2	Compile the interactive application . . . . .	11
9.2.3	Compile the test executable . . . . .	11

# 1 Project architecture (high-level overview)

Our codebase is organized as follows (rooted in `src/`):

```

src/
    main.cpp                  (interactive application)
    tests/
        TestPricing.cpp      (automated consistency tests)
    model/
        BlackScholesModel.hpp
    grid/
        FdGrid.hpp
        GridParameters.hpp
    products/
        InterfaceProducts.hpp
        EuropeanCall.hpp     EuropeanPut.hpp
        AmericanCall.hpp    AmericanPut.hpp
        Future.hpp
        BullCallSpread.hpp   BearPutSpread.hpp
        Straddle.hpp
    solvers/
        ExplicitFdSolver.hpp
        Solver.cpp           (implementation of ExplicitFdSolver::price)

```

## 1.1 How modules interact

The core pricing pipeline is always the same:

1. We define market parameters in `model/BlackScholesModel.hpp`.
2. We instantiate a product (payoff + boundary conditions) from `products/`.
3. We build a finite-difference grid using `grid/GridParameters.hpp`, which returns an `FdGrid` from `grid/FdGrid.hpp`.
4. We call `ExplicitFdSolver::price` (declared in `solvers/ExplicitFdSolver.hpp`, implemented in `solvers/Solver.cpp`).
5. The solver returns price + Greeks (Delta, Gamma) and the entire solution slice at  $t = 0$ .

## 1.2 What the program takes as inputs

At runtime (interactive application in `src/main.cpp`), the user provides:

- Market inputs:  $S_0$ ,  $r$ ,  $\sigma$ ,  $q$ .
- Contract inputs: maturity  $T$ , strike(s) depending on product.
- Numerical input: a spatial resolution parameter `rel_dS` such that  $\Delta S = \text{rel\_dS} \cdot S_0$ .

The test program (`src/tests/TestPricing.cpp`) uses fixed parameter sets and verifies financial identities and numerical properties automatically.

## 2 Model: Black–Scholes–Merton with dividends

### 2.1 Risk-neutral dynamics

We assume the underlying price  $S_t$  follows (under the risk-neutral measure):

$$dS_t = (r - q)S_t dt + \sigma S_t dW_t,$$

where  $r$  is the risk-free rate,  $\sigma$  the volatility, and  $q$  the continuous dividend yield.

### 2.2 Pricing PDE

For a derivative value  $V(t, S)$ , the Black–Scholes–Merton PDE is:

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + (r - q)S \frac{\partial V}{\partial S} - rV = 0.$$

The PDE is solved backward in time with:

- terminal condition (payoff):  $V(T, S) = \Phi(S)$ ,
- boundary conditions as  $S \rightarrow S_{\min}$  and  $S \rightarrow S_{\max}$ .

### 2.3 Implementation in our code

The model parameters are stored in `model/BlackScholesModel.hpp`:

$$(r, \sigma, q),$$

and used both:

- in `solvers/Solver.cpp` to build finite-difference coefficients (drift  $r - q$ , discount rate  $r$ ),
- in product boundary conditions (e.g. call boundary uses  $e^{-q(T-t)}$  and  $e^{-r(T-t)}$ ).

## 3 Solver: explicit finite differences (`solvers/ExplicitFdSolver.hpp`, `solvers/Solver.cpp`)

### 3.1 Terminal condition

For each product, we initialize:

$$V_i^{N_t} = \Phi(S_i),$$

where  $\Phi$  is the payoff, implemented as `option.payoff(S[i])`.

### 3.2 Explicit backward time stepping

Let  $V_i^n \approx V(t_n, S_i)$ . For each time step  $n = N_t - 1, \dots, 0$ , we compute for interior nodes  $i = 1, \dots, N_s - 1$ :

$$V_i^n = A_i V_{i-1}^{n+1} + B_i V_i^{n+1} + C_i V_{i+1}^{n+1},$$

with:

$$\begin{aligned} A_i &= \frac{\Delta t}{2} \left( \frac{\sigma^2 S_i^2}{\Delta S^2} - \frac{(r - q)S_i}{\Delta S} \right), \\ B_i &= 1 - \Delta t \left( \frac{\sigma^2 S_i^2}{\Delta S^2} + r \right), \\ C_i &= \frac{\Delta t}{2} \left( \frac{\sigma^2 S_i^2}{\Delta S^2} + \frac{(r - q)S_i}{\Delta S} \right). \end{aligned}$$

This is implemented in `solvers/Solver.cpp` with the variables A, B, C.

### 3.3 Boundary conditions

At each time layer  $t_n$ , boundaries are imposed through the product interface:

$$V_0^n = \text{option.leftBoundary}(t_n), \quad V_{N_s}^n = \text{option.rightBoundary}(t_n, S_{\max}).$$

This design makes it straightforward to add products: each product specifies its asymptotic behaviour.

### 3.4 American early exercise

For American options, after computing the continuation value, we project:

$$V_i^n \leftarrow \max(V_i^n, \Phi(S_i)),$$

implemented via: `if(option.isAmerican()) Vnew[i]=max(Vnew[i], option.earlyExerciseValue(Si));`. This is the standard approach for PDE-based American pricing.

### 3.5 Outputs of the solver

The solver returns a `Result` struct (see `solvers/ExplicitFdSolver.hpp`) containing:

- `V0`: the full vector  $\{V(0, S_i)\}_{i=0}^{N_s}$ ,
- `price`: interpolated  $V(0, S_0)$ ,
- `delta`, `gamma`: computed at  $S_0$  using finite differences.

## 4 Implemented products (detailed)

Before discussing grid automation and numerical tuning, we present the set of derivatives supported by our engine. Each product implements the common interface in `products/InterfaceProducts.hpp`:

- `maturity()` and `strike()` (or a representative strike),
- `payoff(S)` (terminal condition),
- `leftBoundary(t)` and `rightBoundary(t, Smax)`,
- optional American hooks: `isAmerican()`, `earlyExerciseValue(S)`.

#### 4.1 European Call (`products/EuropeanCall.hpp`)

**Inputs:** strike  $K$ , maturity  $T$ , model  $(r, \sigma, q)$ .

**Payoff:**

$$\Phi(S) = \max(S - K, 0).$$

**Boundaries:**

$$V(t, 0) = 0, \quad V(t, S_{\max}) \approx S_{\max}e^{-q(T-t)} - Ke^{-r(T-t)}.$$

#### 4.2 European Put (`products/EuropeanPut.hpp`)

**Inputs:**  $K, T, (r, \sigma, q)$ .

**Payoff:**

$$\Phi(S) = \max(K - S, 0).$$

**Boundaries:**

$$V(t, 0) \approx Ke^{-r(T-t)}, \quad V(t, S_{\max}) \approx 0.$$

#### 4.3 American Put (`products/AmericanPut.hpp`)

**Payoff:**  $\max(K - S, 0)$ .

**Early exercise:** enabled via `isAmerican()=true`, with:

$$V(t, S) \geq \max(K - S, 0).$$

**Left boundary choice:** we set  $V(t, 0) = K$ . This reflects the fact that when  $S = 0$ , immediate exercise yields  $K$  and is optimal.

#### 4.4 American Call (`products/AmericanCall.hpp`)

**Payoff:**  $\max(S - K, 0)$ .

**Early exercise:** enabled. In Black–Scholes, early exercise for calls is:

- never optimal when  $q = 0$  (American call equals European call),
- potentially optimal when  $q > 0$  (dividends can make early exercise attractive).

**Boundary at  $S_{\max}$ :** same asymptotic as European call:

$$V(t, S_{\max}) \approx S_{\max}e^{-q(T-t)} - Ke^{-r(T-t)}.$$

#### 4.5 Forward/Future-like contract (`products/Future.hpp`)

**Inputs:** delivery price  $K$ , maturity  $T$ .

**Payoff:**

$$\Phi(S) = S - K.$$

This product admits a closed-form price:

$$V(0, S_0) = S_0e^{-qT} - Ke^{-rT}.$$

**Boundaries:**

$$V(t, 0) \approx -Ke^{-r(T-t)}, \quad V(t, S_{\max}) \approx S_{\max}e^{-q(T-t)} - Ke^{-r(T-t)}.$$

## 4.6 Bull Call Spread (products/BullCallSpread.hpp)

**Inputs:** strikes  $K_1 < K_2$ , maturity  $T$ .

**Payoff:**

$$\Phi(S) = \max(S - K_1, 0) - \max(S - K_2, 0).$$

This payoff is bounded between 0 and  $(K_2 - K_1)$ .

**Right boundary:** for very large  $S$ , the payoff saturates to  $K_2 - K_1$ , hence:

$$V(t, S_{\max}) \approx (K_2 - K_1)e^{-r(T-t)}.$$

**Left boundary:**  $V(t, 0) = 0$ .

## 4.7 Bear Put Spread (products/BearPutSpread.hpp)

**Inputs:**  $K_1 < K_2$ ,  $T$ .

**Payoff:**

$$\Phi(S) = \max(K_2 - S, 0) - \max(K_1 - S, 0),$$

also bounded between 0 and  $(K_2 - K_1)$ .

**Left boundary:** for  $S = 0$ , payoff saturates to  $K_2 - K_1$ , so:

$$V(t, 0) \approx (K_2 - K_1)e^{-r(T-t)}.$$

**Right boundary:**  $V(t, S_{\max}) \approx 0$ .

## 4.8 Straddle (products/Straddle.hpp)

**Inputs:** strike  $K$ , maturity  $T$ .

**Payoff:**

$$\Phi(S) = |S - K| = \max(S - K, 0) + \max(K - S, 0).$$

A key identity is:

$$\text{Straddle} = \text{Call}(K) + \text{Put}(K).$$

**Boundaries:**

$$V(t, 0) \approx Ke^{-r(T-t)}, \quad V(t, S_{\max}) \approx S_{\max}e^{-q(T-t)} - Ke^{-r(T-t)}.$$

## 5 Finite-difference grid and interpolation

### 5.1 Grid definition (grid/FdGrid.hpp)

We discretize:

$$t_n = n\Delta t, \quad n = 0, \dots, N_t, \quad S_i = S_{\min} + i\Delta S, \quad i = 0, \dots, N_s.$$

The **FdGrid** class stores:

$$T, S_{\min}, S_{\max}, N_t, N_s, \Delta t, \Delta S,$$

and exposes **timeGrid()**, **priceGrid()**, plus a safe linear interpolation routine **interpolate(V, S0)** to evaluate  $V(0, S_0)$  even when  $S_0$  is not exactly a grid node.

## 5.2 Interpolation used for the final price

After solving for the vector  $V(0, S_i)$  on the grid, we compute:

$$V(0, S_0) \approx (1 - w)V(0, S_i) + wV(0, S_{i+1}), \quad w = \frac{S_0 - S_i}{S_{i+1} - S_i},$$

with  $(S_i, S_{i+1})$  such that  $S_i \leq S_0 \leq S_{i+1}$ . This is implemented in `FdGrid::interpolate`.

## 6 Automatic grid parameter selection: stability vs accuracy vs speed

A key engineering challenge of this project was to select grid parameters that simultaneously ensure:

- **stability** (explicit scheme  $\Rightarrow$  CFL-type constraint),
- **accuracy** (small discretization error near  $S_0$  and around the strike),
- **efficiency** (runtime and memory remain reasonable for interactive use).

In practice, this was the main difficulty: a naive grid choice can easily produce a solver that is either (i) unstable and diverges, or (ii) stable but too slow, or (iii) fast but inaccurate because the spatial mesh becomes too coarse.

### 6.1 Domain truncation: choosing $S_{\max}$ via a lognormal quantile

Under Black–Scholes,  $\log S_T$  is Gaussian:

$$\log S_T \sim \mathcal{N}\left(\log S_0 + (r - q - \frac{1}{2}\sigma^2)T, \sigma^2 T\right).$$

We choose the upper boundary as a high quantile of the lognormal distribution:

$$S_{\max} = S_0 \exp\left((r - q - \frac{1}{2}\sigma^2)T + z\sigma\sqrt{T}\right),$$

with a fixed safety level  $z$  (e.g.  $z = 5$ ). The motivation is to make the probability of  $S_T > S_{\max}$  negligible, so that boundary conditions do not contaminate the solution in the region of interest (near  $S_0$ ). We set  $S_{\min} = 0$  for all products.

### 6.2 Space resolution: choosing $N_s$ from a relative mesh size

A first implementation used fixed values of  $N_s$  (e.g. 250/400/800). This worked on some parameter sets but degraded when  $S_{\max}$  increased (large  $T$  or large  $\sigma$ ). Since

$$\Delta S = \frac{S_{\max} - S_{\min}}{N_s},$$

a larger  $S_{\max}$  mechanically increases  $\Delta S$  and reduces accuracy around  $S_0$ .

To fix this, our final design makes the spatial step proportional to the spot:

$$\Delta S = \text{rel\_dS} \cdot S_0,$$

where `rel_dS` is chosen by the user at runtime (typical values: 0.004 fast, 0.002 balanced, 0.001 accurate). We then set:

$$N_s = \left\lceil \frac{S_{\max} - S_{\min}}{\Delta S} \right\rceil.$$

This ensures that local resolution near  $S_0$  remains comparable across parameter sets, even when  $S_{\max}$  varies substantially.

### 6.3 Time resolution: explicit stability constraint and its implications

The explicit finite-difference scheme is only conditionally stable, so  $\Delta t$  must satisfy a CFL-type bound. A conservative practical constraint is:

$$\Delta t \lesssim \frac{1}{\sigma^2 S_{\max}^2 / \Delta S^2 + r}.$$

We choose  $\Delta t$  using this worst-case bound (with a safety factor) and set:

$$N_t = \left\lceil \frac{T}{\Delta t} \right\rceil.$$

### 6.4 Main criticism of the approach

The above procedure works reliably for many cases, but it also highlights the main limitation of our numerical method:

- Because we use an **explicit** scheme, stability can force  $N_t$  to become very large when  $\sigma$  or  $S_{\max}$  is large, which may make the solver slow.
- Using a uniform grid in  $S$  can still be inefficient: accuracy is needed mainly near  $S_0$  (and around strikes), while far-away regions consume grid points but contribute little to the price at  $S_0$ .

These limitations are well-known in PDE pricing: implicit schemes (Backward Euler / Crank–Nicolson) remove the strict stability constraint, and non-uniform grids (or a change of variable to  $\log S$ ) concentrate points where they matter most. We chose the explicit scheme for simplicity and transparency, but grid automation was crucial to obtain a solver that is both reasonably fast and accurate in practice.

## 7 Greeks: formulas and numerical computation

### 7.1 Definitions (continuous-time)

For a derivative price  $V(0, S_0; r, \sigma, q, T, \dots)$ , the standard Greeks are:

$$\begin{aligned} \Delta &= \frac{\partial V}{\partial S_0}, & \Gamma &= \frac{\partial^2 V}{\partial S_0^2}, & \Theta &= \frac{\partial V}{\partial t}, \\ \nu \text{ (Vega)} &= \frac{\partial V}{\partial \sigma}, & \rho &= \frac{\partial V}{\partial r}. \end{aligned}$$

In this project, we compute **Delta** and **Gamma** directly from the PDE solution on the  $S$ -grid. The other Greeks (Theta, Vega, Rho) can be computed by bump-and-revalue, but we chose not to implement them to keep the engine focused.

## 7.2 Discrete formulas used in our code

Let  $V_i \approx V(0, S_i)$  and choose  $i^*$  such that  $S_{i^*}$  is the closest grid node to  $S_0$  (while avoiding boundaries). Then:

$$\Delta(0, S_0) \approx \frac{V_{i^*+1} - V_{i^*-1}}{2\Delta S},$$

$$\Gamma(0, S_0) \approx \frac{V_{i^*+1} - 2V_{i^*} + V_{i^*-1}}{\Delta S^2}.$$

We avoid  $i^* = 0$  and  $i^* = N_s$  to prevent unreliable boundary derivatives; this is implemented in `solvers/Solver.cpp` by shifting the index away from the boundaries.

## 8 Test suite: consistency checks

The test executable `bs_tests` runs a suite of checks comparing our numerical outputs to basic financial identities and numerical properties:

- Put–call parity (European):

$$C - P = S_0 e^{-qT} - K e^{-rT}.$$

- Forward closed-form price:

$$V_{\text{fwd}}(0) = S_0 e^{-qT} - K e^{-rT}.$$

- American dominance:  $P^{Am} \geq P^{Eu}$  and  $C^{Am} \geq C^{Eu}$ .
- Straddle consistency: Straddle  $\approx C + P$ .
- Spread bounds: spreads are bounded by  $(K_2 - K_1)e^{-rT}$ .
- Greek sanity checks: for forwards,  $\Delta \approx e^{-qT}$  and  $\Gamma \approx 0$ .

These tests were crucial to ensure our code behaves correctly and to detect regressions after refactoring.

## 9 How to build and run the code

The full source code is publicly available on GitHub:

[https://github.com/Samgit0532/Pricing\\_par\\_EDP](https://github.com/Samgit0532/Pricing_par_EDP).

This section briefly explains how to compile and run the project, either using `CMake` or directly with `g++`. A `README.md` is also there for more details.

### 9.1 Build and run with CMake

#### 9.1.1 Requirements

- `CMake` version  $\geq 3.16$ ,
- a C++ compiler supporting the `C++17` standard (e.g. `g++`).

### 9.1.2 Build

From the root directory of the project, run:

```
cmake -S . -B build
cmake --build build -j
```

This generates two executables in the `build/` directory.

### 9.1.3 Run the interactive application

```
./build/bs_app
```

The program launches an interactive interface where the user selects:

- the financial product,
- market parameters ( $S_0, r, \sigma, q$ ),
- contract parameters (maturity, strike(s)),
- the numerical resolution parameter `rel_dS`.

### 9.1.4 Run the test suite

```
./build/bs_tests
```

The test executable automatically checks pricing identities and numerical sanity properties (put–call parity, forward pricing, American dominance, bounded spreads, and Greek consistency).

## 9.2 Build and run without CMake (direct compilation)

### 9.2.1 Requirements

- g++ with C++17 support.

### 9.2.2 Compile the interactive application

```
g++ -std=c++17 -O2 -O2 -I./src src/main.cpp src/solvers/Solver.cpp -o bs_app
```

Run:

```
./bs_app
```

### 9.2.3 Compile the test executable

```
g++ -std=c++17 -O2 -I./src src/tests/TestPricing.cpp src/solvers/Solver.cpp -o bs_tests
```

Run:

```
./bs_tests
```